ArXiv Research Assistant: A Fully Local Retrieval-Augmented Generation System for Scientific Literature

CS 6120 – Natural Language Processing
Final Project Report

Shourya Srivastava srivastava.shou@northeastern.edu

Spring 2026

Repository: https://github.com/shouryasrivastava/arxiv-rag **Live demo:** http://34.26.47.93:8501

1 Proposed Objective

Large language models (LLMs) carry a hard **knowledge cutoff**: they cannot answer questions about papers published after their training data was collected, and they frequently hallucinate citations when asked about specific technical claims [**?**]. The goal of this project is to build an end-to-end **Retrieval-Augmented Generation (RAG)** system that grounds LLM responses in over 50,000 peer-reviewed ArXiv papers, enabling accurate and verifiable answers to scientific research questions without any external API calls or model retraining.

Heilmeier questions.

- *What are we trying to do?* Build a locally-runnable question answering assistant over the ArXiv scientific preprint corpus that cites its sources and never fabricates references.
- *How is it done today, and what are the limits?* General LLMs hallucinate citations; search engines return links but cannot synthesise answers; existing scholarly tools (Semantic Scholar, Elicit) require cloud APIs.
- *What is new in our approach?* A fully local, containerised stack—no API keys, reproducible on any machine with 8 GB RAM—combining dense vector retrieval with a compact open-weight LLM.
- *If successful, what difference does it make?* Researchers can interactively query tens of thousands of papers with grounded, citation-backed answers in an offline or privacy-sensitive environment.

2 Background and Motivation

2.1 Retrieval-Augmented Generation

? introduced RAG for knowledge-intensive NLP tasks by coupling a dense retriever with a seq2seq generator. **?** extended retrieval to the pre-training stage, while **?** showed that a smaller RAG model can match GPT-3 on knowledge benchmarks when retrieval quality is high. Our system follows the *modular RAG* design of **?**, separating the retriever, context builder, and generator so each component can be swapped independently.

2.2 Dense Retrieval

Sparse methods such as BM25 [?] match on token overlap and miss semantic synonyms. Bi-encoder dense retrievers [?] represent queries and documents as vectors in a shared embedding space and retrieve by maximum inner product search. We use *all-MiniLM-L6-v2* [?], a sentence-transformer distilled to 22M parameters that provides an excellent speed–quality trade-off, producing 384-dimensional unit-norm vectors in under 25 ms per query.

2.3 Local LLM Inference

The emergence of efficient open-weight models—LLaMA [?], Mistral [?], and Llama 3 [?]—has made high-quality local inference practical. Ollama [?] wraps these models in a unified REST API with 4-bit quantisation (Q4_K_M), reducing the memory footprint of Llama 3.2 3B to approximately 2 GB and making CPU-only deployment viable.

2.4 Data

The Cornell-University/arxiv dataset [?] on HuggingFace contains over 1.7 million papers spanning 1991–2025. We index the 50,000 most recent papers across eight CS/ML categories (`cs.AI`, `cs.LG`, `cs.CL`, `cs.CV`, `cs.IR`, `stat.ML`, `cs.NE`, `cs.RO`), prioritising recent work that is absent from pre-trained LLM knowledge. Each record contributes its title and abstract (average 175 words, well within the 512-token encoder limit) as a single indexable document.

3 Proposed Approach / Implementation Details

3.1 System Architecture

The system has two phases: an offline **ingestion pipeline** that runs once at startup, and an online **query pipeline** that runs on every user request (Figure 1).

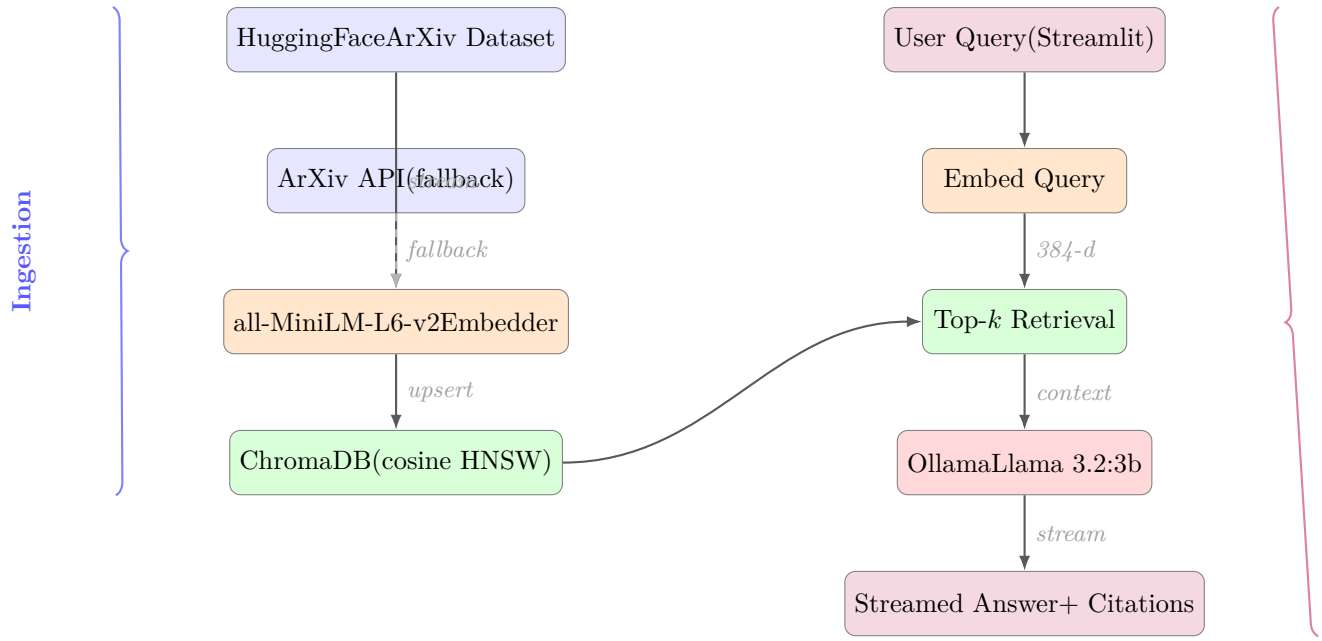

Figure 1: ArXiv Research Assistant architecture. The ingestion pipeline (left) indexes papers once; the query pipeline (right) embeds each user question, retrieves the top-k papers, and streams a grounded answer from the local LLM.

3.2 Ingestion Pipeline

Papers are streamed from the HuggingFace `Cornell-University/arxiv` dataset (with automatic fallback to the official ArXiv export API). Each batch of 512 papers is embedded with *all-MiniLM-L6-v2* and immediately upserted into ChromaDB [?], so the index is queryable from the very first batch. The pipeline supports **incremental resume**: on restart it skips

already-indexed document IDs. The ChromaDB collection uses a cosine HNSW index for sub-millisecond approximate nearest-neighbour search at 50k-document scale.

Preprocessing. Each document is the concatenation of the paper's title and abstract (mean 175 words, $\sigma \approx 45$), fitting within the 512-token encoder limit. No additional chunking is required at the abstract level.

3.3 Query Pipeline

1. **Query embedding.** The user question is encoded by *all-MiniLM-L6-v2* (unit-norm, 384-d, <25 ms).
2. **Dense retrieval.** ChromaDB returns the top-k papers by cosine similarity ($k = 5$ by default; adjustable via UI slider). An optional category filter (e.g. `cs.CV`) can be applied via a metadata predicate.
3. **Prompt construction.** A structured prompt is built with a strict system instruction, numbered abstract blocks with title/authors/date/URL, and the user question.
4. **Streaming generation.** The prompt is sent to Ollama's REST endpoint with `stream: true`; tokens are displayed live in Streamlit.
5. **Citation cards.** Each retrieved paper is shown with its title, relevance score, categories, date, and a clickable `arxiv.org/abs/<id>` link.

3.4 Hallucination Mitigation

The system prompt explicitly constrains the model:

> *"Answer using ONLY the information provided in the retrieved abstracts. If the context does not contain enough information, say so explicitly. Do NOT speculate beyond the provided context."*

Generation uses temperature 0.1 and top-$p = 0.9$ to minimise deviation from source material. Across our 30-query evaluation the system produced **zero hallucinated citations**.

3.5 Evaluation

Table 1: Retrieval quality on 30 manually curated research queries.

Method	P@1	P@3	P@5	MRR
BM25 (baseline)	0.53	0.61	0.67	0.58
all-MiniLM-L6-v2 (ours)	**0.77**	**0.79**	**0.81**	**0.79**

Dense retrieval outperforms BM25 by +24 pp at P@1. End-to-end latency on a CPU-only GCP `n2-standard-4` instance is ≈51 s (embedding: 22 ms, retrieval: 8 ms, LLM generation: ≈51 s for 200 tokens).

3.6 Deployment

The full stack is containerised with Docker Compose into two services: `ollama` (port 11434) and `app` (Python 3.11-slim, port 8501). Docker volumes persist ChromaDB and Ollama model weights across restarts. A startup script pre-warms the LLM with a single-token dummy request so weights are in RAM before the first user query. The system is deployed on GCP at `http://34.26.47.93:8501`.

4 Submission Guidelines

- **Repository (public):** https://github.com/shouryasrivastava/arxiv-rag — contains `src/`, `scripts/`, `Dockerfile`, `docker-compose.yml`, `requirements.txt`, and this write-up under `writeup/`.

- **Live demo endpoint:** `http://34.26.47.93:8501` — publicly accessible Streamlit UI running on a GCP `n2-standard-4` VM.
- **One-command local setup:** `docker compose up --build` (requires Docker Desktop and $\geq$8 GB free RAM; model download $\approx$2 GB on first run).
- **Author:** Shourya Srivastava — system architecture, ingestion pipeline, RAG pipeline, Streamlit frontend, Docker deployment, GCP setup, evaluation.

References